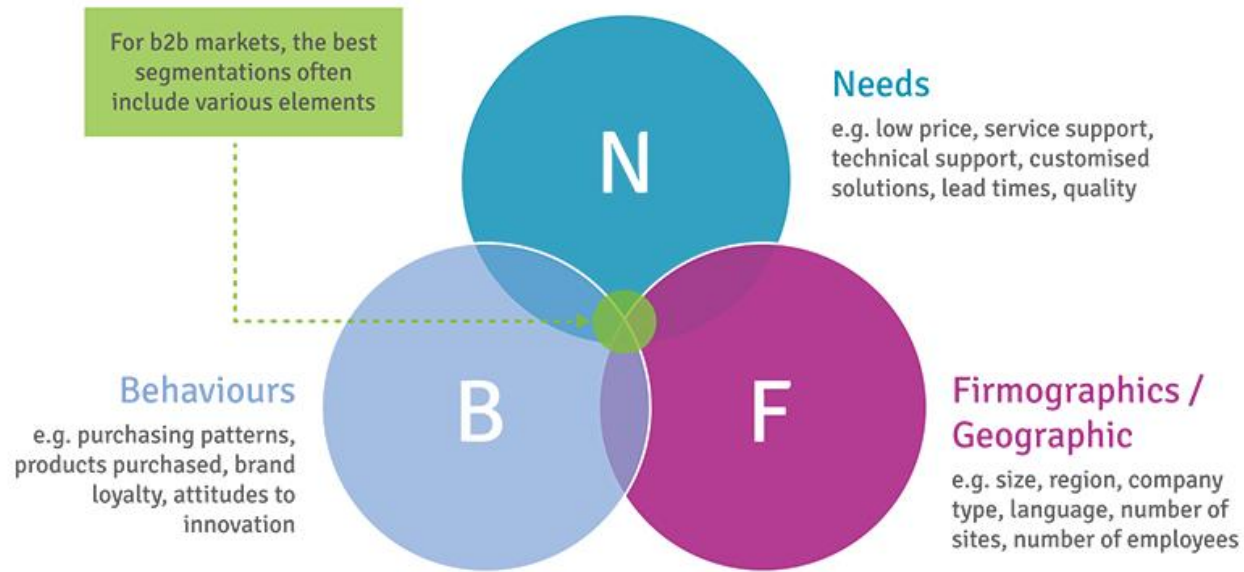


Hierarchical Clustering for B2B Customer Segmentation





Problem Statement



Artifact
Submission



Recommendation

Data Description



Project Evaluation
Criteria



Vibelines
& Bounties



1. Understanding Hierarchical Clustering - The Basics

Hierarchical Clustering is an unsupervised machine learning algorithm that builds a hierarchy of clusters, rather than requiring a pre-defined number of clusters (like K-Means). It creates a tree-like structure called a **dendrogram**, which visually represents the nested relationships between clusters.

There are two main types of hierarchical clustering:

- **Agglomerative (Bottom-Up):** This is the most common approach.
 - Starts with each data point as its own individual cluster.
 - Iteratively merges the closest pairs of clusters until all data points are in a single cluster, or a stopping criterion is met.
- **Divisive (Top-Down):**
 - Starts with all data points in one large cluster.
 - Recursively splits the clusters into smaller clusters until each data point is in its own cluster, or a stopping criterion is met.

The key advantage of hierarchical clustering is that it doesn't require you to specify the number of clusters (K) upfront. Instead, you can decide on the number of clusters by visually inspecting the dendrogram or by using a specific threshold.



2. Associated Concepts in Hierarchical Clustering

Hierarchical clustering relies on several key concepts and considerations:

- **Unsupervised Learning:** Like K-Means, hierarchical clustering is an unsupervised algorithm. It discovers patterns or groupings within the data without any prior knowledge of what those groups should be.
- **Distance Metric (Proximity Measure):** This defines how the "closeness" or "similarity" between individual data points is measured. Common metrics include:
 - **Euclidean Distance:** The straight-line distance between two points in a multi-dimensional space.
 - **Manhattan Distance:** The sum of the absolute differences of their Cartesian coordinates.
- **Linkage Criterion:** This defines how the "distance" between two *clusters* (not just individual points) is calculated. This is crucial for determining which clusters to merge (in agglomerative) or split (in divisive). Common linkage methods include:
 - **Single Linkage:** The distance between the closest points in the two clusters. (Can lead to "chaining" effect).
 - **Complete Linkage:** The distance between the farthest points in the two clusters. (Tends to produce more compact, spherical clusters).
 - **Average Linkage:** The average distance between all pairs of points in the two clusters.
 - **Ward's Method:** Minimizes the variance within each cluster when merging. (Often produces good, balanced clusters).
- **Dendrogram:** This is the primary output of hierarchical clustering. It's a tree-like diagram that illustrates the sequence of merges or splits.
 - The height of the merge point in a dendrogram indicates the distance (or dissimilarity) between the clusters being merged.
 - You can "cut" the dendrogram at a certain height to obtain a desired number of clusters.
- **Feature Scaling:** It is **essential** to scale your features (e.g., using StandardScaler to achieve zero mean and unit variance) before applying hierarchical clustering. This is because it is a distance-based algorithm, and features with larger numerical ranges would disproportionately influence the distance calculations, leading to biased clustering.
- **Cluster Profiling:** After determining the clusters, it's crucial to analyze the characteristics (e.g., average feature values, distributions) of the data points within each cluster to understand what defines that segment.



3. Why Hierarchical Clustering is Important and in What Industries

Hierarchical clustering is a powerful technique for segmenting data, particularly when the underlying structure of clusters is unknown or when a visual hierarchy is beneficial.

Why is Hierarchical Clustering Important?

- **No Need for Predefined 'K':** Unlike K-Means, you don't need to specify the number of clusters beforehand. This is a significant advantage when you have no prior intuition about the optimal number of groups.
- **Visual Interpretation with Dendrograms:** The dendrogram provides a clear, intuitive visualization of how clusters are formed and their relationships, allowing for flexible cluster selection.
- **Reveals Nested Structures:** Can uncover sub-clusters within larger clusters, providing a more granular understanding of the data.
- **Flexible Cluster Granularity:** You can choose the level of granularity for your clusters by cutting the dendrogram at different heights.
- **Customer Segmentation:** Identifies distinct groups of customers with similar behaviors, preferences, or demographics, enabling targeted marketing and personalized experiences.
- **Market Research:** Uncovers natural groupings within survey responses or consumer data to understand market segments.
- **Biology & Genomics:** Grouping similar species, genes, or proteins based on their characteristics.



4. Industries where Hierarchical Clustering is particularly useful:

- **Wholesale & Distribution (Core Application):** Segmenting business clients based on purchasing volume, product categories, and distribution channels.
- **Biology & Bioinformatics:** Classifying species, genes, or proteins based on genetic or phenotypic similarities.
- **Market Research:** Understanding consumer segments from survey data or behavioral patterns, especially when exploring new markets.
- **Customer Relationship Management (CRM):** Segmenting customers for personalized marketing, product recommendations, and loyalty programs.
- **Social Sciences:** Grouping individuals based on survey responses, attitudes, or behaviors.
- **Document Analysis:** Clustering similar documents or articles based on their content.
- **Image Processing:** Grouping similar image regions or objects.
- **Logistics & Supply Chain:** Segmenting clients based on shipping volume, delivery requirements, and geographical location.



Data Description

This project focuses on applying **Hierarchical Clustering** to a dataset containing annual spending patterns of different business clients across various product categories. The objective is to identify distinct segments of B2B customers and to visualize the hierarchical relationships between these segments using a dendrogram. This approach will enable the business to tailor its sales, marketing, and distribution strategies more effectively without needing to pre-define the number of client groups.

Dataset Details:

- **Dataset Name:** Annual spends by different clients in respective product category

Column description (Key Features for Clustering):

1. **Channel:** Distribution channel type (e.g., Horeca, Retail). This is a categorical feature that might need to be encoded or used for initial filtering.
2. **Region:** Region code (location of the client). Categorical, also potentially useful for initial filtering or as a characteristic of segments.
3. **Fresh:** Spend by clients in the Fresh Category.
4. **Milk:** Spend by clients in the Milk Category.
5. **Grocery:** Spend by clients in the Grocery Category.
6. **Frozen:** Spend by clients in the Frozen Category.
7. **Detergents_Paper:** Spend by clients in the Detergents paper Category.
8. **Delicassen:** Spend by clients in the Delicassen Category.



Artifact Submission

Your submission must include the following five artifacts, all packaged within a single GitHub repository.

1. Jupyter Notebook (.ipynb) This is the core of your submission. Your Jupyter Notebook should be a complete, well-documented narrative of your data analysis journey. It must include:

- **Detailed Explanations:** Use Markdown cells to explain your thought process, the questions you are trying to answer, and the insights you've uncovered.
- **Clean Code:** The code should be well-structured, easy to read, and free of unnecessary clutter.
- **Comprehensive Comments:** Use comments to explain complex logic and the purpose of different code blocks.
- **Key Visualizations:** All visualizations should be clear, properly labeled, and directly support your findings.

2. Presentation (.pptx or .pdf)

Create a compelling presentation that summarizes your team's analysis and key findings. This presentation should serve as your final pitch. It must include:

- **Executive Summary:** A concise overview of your findings.
- **Key Insights:** The most important takeaways from your analysis.
- **Data-Driven Recommendations:** Actionable steps that can be taken based on your insights.
- **Supporting Visualizations:** A selection of your best visualizations to illustrate your points.

3. README File (.md)

The README file is the first thing we'll look at. It should serve as a quick guide to your project and provide essential details. It must include:

- **Project Title :**
- **Brief Problem Statement:** A summary of the project and your approach.
- **Summary of Findings:** A bullet-point summary of your most significant insights.

4. Attached Dataset

Please include the original dataset (.csv or other format) within your repository. This ensures the judges can reproduce your analysis without any issues.

5. GitHub Repository

Your final submission will be your GitHub repository. The repository name **must follow this exact format:**

Clustering_ProjectName_TMP



Challenge Evaluation Criteria

Criteria Name	Criteria weight
Data Understanding and Exploratory Data Analysis	20%
Data preprocessing and feature engineering	25%
Model building and evaluation	30%
Business Recommendation	15%
Coding guidelines and standards	10%



Recommendation for Hierarchical Clustering

- **Data Preprocessing:**
 - Selecting the numerical columns representing annual spending (Fresh, Milk, Grocery, Frozen, Detergents_Paper, Delicassen).
 - **Handling categorical features:** Decide whether to include Channel and Region in the clustering. If included, they will need to be one-hot encoded. Alternatively, clustering can be performed on spending data first, and then segments can be profiled by their dominant Channel and Region.
 - **Crucially, performing feature scaling** on the spending columns (e.g., using StandardScaler). This is essential because hierarchical clustering is a distance-based algorithm, and features with larger numerical ranges would disproportionately influence the distance calculations if not scaled.
- **Distance Matrix Calculation:**
 - Calculating the pairwise distances between all client data points using a chosen distance metric (e.g., Euclidean distance).
- **Linkage Method Application:**
 - Applying a chosen linkage criterion (e.g., Ward's method, Average linkage) to define the distance between clusters and build the hierarchy. Ward's method is often a good starting point for numerical data as it tends to produce compact clusters.
- **Dendrogram Visualization:**
 - Generating and visualizing the **dendrogram**. This tree-like diagram will show how individual clients are successively merged into larger clusters based on their spending similarities.
- **Determining the Number of Clusters:**
 - By visually inspecting the dendrogram, identifying a suitable "cut-off" point (a horizontal line) that yields a meaningful number of clusters. This allows for flexibility in choosing the granularity of segmentation, for example, distinguishing between "Large Grocery Chains," "Small Independent Restaurants," or "Specialty Food Retailers."
- **Cluster Assignment & Profiling:**
 - Assigning each client to a specific cluster based on the chosen cut-off point on the dendrogram.
 - Analyzing the characteristics (e.g., average spend in each category, dominant Channel and Region) of clients within each identified cluster. For example, one cluster might be "High-Volume Fresh & Milk Distributors to Horeca in Region A," while another could be "Grocery & Detergent Focused Retailers in Region B."



Project Outcomes

The outcome of this project will be a clear, visually interpretable segmentation of B2B customers based on their annual spending patterns across different product categories, along with an understanding of the hierarchical relationships between these groups. This insight can be invaluable for:

- **Sales Strategy:** Tailoring sales pitches and product recommendations to the specific needs and purchasing habits of each client segment.
- **Marketing Campaigns:** Designing targeted promotions or loyalty programs that resonate with the spending profiles of different client groups.
- **Distribution Optimization:** Understanding which channels and regions are dominant for certain client segments.
- **Account Management:** Prioritizing and allocating resources to different client segments based on their value and potential.
- **Inventory Planning:** Forecasting demand more accurately by understanding the purchasing patterns of client segments.



Lets Go

