

## Building a Netflix recommender system leveraging Kmeans Minibatch

K-Means is a popular clustering algorithm that aims to partition 'n' observations into 'k' clusters, where each observation belongs to the cluster with the nearest mean (centroid). The standard K-Means algorithm processes the entire dataset in each iteration to update the cluster centroids. This can be computationally expensive and slow for very large datasets.



Mini-Batch K-Means is a variation of the K-Means algorithm designed to address this scalability issue. Instead of using the full dataset, Mini-Batch K-Means performs updates based on small, randomly sampled subsets of the data called "mini-batches".

Here's how it works:

- **Initialize Centroids:** Just like standard K-Means, it starts by randomly initializing 'k' cluster centroids.

Iterative Updates with Mini-Batches:

- In each iteration, a small, random subset of data (a mini-batch) is drawn from the dataset.
- For each point in the mini-batch, its closest centroid is identified. The chosen centroids are then updated incrementally using a weighted average of the new data points assigned to them. This update is often a learning rate-like adjustment, rather than a full re-calculation based on all points assigned to that centroid in the entire dataset.
- **Convergence:** The process continues for a fixed number of iterations or until the centroids stabilize.

## Key Advantages of Mini-Batch K-Means:

- **Speed for Large Datasets:** Significantly faster than standard K-Means for very large datasets because it doesn't need to load and process the entire dataset in memory for every update.
- **Memory Efficiency:** Requires less memory as it only operates on small mini-batches at a time.
- **Approximation:** While it's an approximation of K-Means, its results are often very close to those of the full K-Batch algorithm, especially with a sufficient number of iterations and well-chosen mini-batch sizes.
- **Using Mini-Batch K-Means in Content-Based Filtering for Netflix**  
Content-based filtering recommends items to users based on the attributes (content) of items they have liked in the past. For Netflix movies and shows, these attributes can include genres, directors, cast, keywords from descriptions, and even plot summaries.

Here's how Mini-Batch K-Means can be leveraged in such a system:

## Feature Extraction and Representation:

- **Movie/Show Content Data:** Start with your Netflix dataset containing features like title, genres, director, cast, description, etc..
- **Text Preprocessing:** Clean and preprocess textual features (like description, genres, director, cast) using techniques like lowercasing, removing stop words, lemmatization, and potentially replacing spaces with underscores for multi-word entities (e.g., "science fiction" -> "science\_fiction").
- **Vectorization (Embeddings):** Convert these textual features into numerical vectors.
- **TF-IDF:** If your features are simple bags of words (e.g., combined all\_features string), you can use TF-IDF to get a numerical representation for each movie/show.
- **Word Embeddings (e.g., Word2Vec, FastText):** A more sophisticated approach is to use pre-trained (or custom-trained) word embeddings. You can average the embeddings of all keywords/genres/directors in a

movie's `all_features` list to get a single "movie embedding" vector. This captures semantic meaning better.

- Node2Vec (as discussed previously): If you build a graph connecting movies to genres, directors, actors, etc., Node2Vec can generate embeddings for each movie node and each feature node. You'd then use these movie embeddings. Result: Each movie/show in your Netflix catalog is now represented as a dense numerical vector. This is your "dataset" for clustering.

#### Clustering Movies/Shows with Mini-Batch K-Means:

- Why Cluster? The idea is that movies/shows within the same cluster are "content-similar" to each other. They share common genres, directors, cast members, or thematic elements.

#### Applying Mini-Batch K-Means:

- You would feed the vectorized movie/show data (e.g., 20-dimensional embeddings from Node2Vec, or 300-dimensional averaged Word2Vec vectors) to the Mini-Batch K-Means algorithm.
- You'll need to choose the number of clusters, 'k'. This can be determined using methods like the elbow method or silhouette score.
- Mini-Batch K-Means will efficiently group your vast Netflix catalog into 'k' content-based clusters.

#### Recommendation Generation:

- User Profile: When a user likes a particular movie or show, you identify which cluster that liked item belongs to.

#### Recommendation Strategy:

- Intra-Cluster Recommendation: Recommend other movies/shows from the same cluster that the user has not yet seen or rated. This is the most straightforward content-based recommendation.
- Inter-Cluster Exploration: You could also recommend items from neighboring clusters (clusters whose centroids are numerically close in the embedding space) to offer a slightly diversified but still relevant set of recommendations.

- **User-Centric Clustering (Optional):** If you have user-specific content preference vectors (e.g., average of embeddings of all movies a user has watched), you could also cluster users based on these vectors using Mini-Batch K-Means. Then, for a given user, recommend items from clusters preferred by users in their cluster.

Benefits of using Mini-Batch K-Means here:

- **Scalability for Netflix Catalog:** Netflix has a massive catalog. Standard K-Means would be too slow and memory-intensive. Mini-Batch K-Means allows you to cluster millions of items efficiently.
- **Fast Model Updates:** If new movies are added frequently, you could potentially update clusters or re-assign new items to existing clusters much faster than with full K-Means.
- **Personalization:** By grouping content, you can easily identify content niches and recommend items that align with a user's demonstrated content preferences, even for new users (cold start) or new items.
- **Complementary to Collaborative Filtering:** Content-based recommendations (from clustering) can be combined with collaborative filtering techniques (which look at user-item interactions) to build a robust hybrid recommender system.